


Static W8 authoring-to-device path

Version-locked source, recipe, AOT target, runtime, and evidence



W8 per-tensor K-means palettization

INT8 per-tensor quantization · FP16 fixed-size · 4,096 tokens